

Danil Korchagin

Data Engineer

📍 Prague, CZ | ✉ kordp51@gmail.com | 📞 +420775237395 | 🔗 www.linkedin.com/in/danil-korchagin-engineer | ⚙️ <https://github.com/Danil5115>

SUMMARY

Data Engineer who ships reliable, observable data products end-to-end. I design lean batch and streaming pipelines (Airflow, Spark, Kafka), model clean schemas and marts, and keep SLAs predictable without drama. I turn vague business questions into BI-ready datasets with clear data contracts and quality gates, and I'm comfortable owning delivery while coordinating analysts, engineers, and stakeholders. Hands-on and pragmatic, I lead small initiatives - from ingestion and orchestration to serving layers and dashboards - so teams get faster answers and fewer surprises.

TECHNICAL SKILLS

-**Data Engineering & Orchestration:** Python, SQL, Apache Airflow, Apache Spark (PySpark, Spark SQL, Structured Streaming), Apache Kafka, dbt, Databricks
-**Databases, Warehousing & Lake:** PostgreSQL, ClickHouse, Vertica, MongoDB, Redis, S3, Parquet, Apache Hadoop
-**Platform, BI & Ops:** Kubernetes, Docker, GitLab CI/CD, Git, Apache Superset, Metabase

PROFESSIONAL EXPERIENCE

Data Engineer

DCUK | Ústí nad Labem, Czechia | Oct 2024 – Present

- Developed complex data transformations with PySpark for scheduled analytics pipelines, reducing compute usage by 23.9% and keeping nightly SLAs stable.
- Implemented dataset-aware scheduling and deferrable sensors in Airflow, eliminating cross-DAG polling and stabilizing dependency-driven pipelines (lower scheduler load, fewer stale runs).
- Delivered a Kafka → Spark Structured Streaming path with S3 checkpoints into a serving layer, enabling near-real-time widgets (median latency ~12–15 min) for priority datasets.
- Redesigned the serving model and marts, speeding up BI queries and removing peak-time contention for critical dashboards.

Tech Stack: Python, SQL, Airflow, Spark (PySpark, Spark SQL, Structured Streaming), Kafka, PostgreSQL, ClickHouse, S3, Kubernetes, Superset, Docker

Data Engineer

Exportio | Prague, Czechia | Apr 2023 – Oct 2024

- Integrated marketplace & back-office data into a Lake-lite → PostgreSQL DWH and automated data prep in Airflow to deliver listing-ready datasets, reducing listing-prep time by 22.6% and cutting pre-publish defects by 26.1%.
- Implemented channel-specific listing automation in Airflow for Amazon, Allegro, Alza, and Kaufland, with automated submission, status tracking, and error triage, which increased first-pass acceptance by +14.8 pp.
- Established production-grade data contracts and quality gates across pipelines with SLAs and live dashboards, which increased DQ pass rate by +11.7 pp.
- Worked with analysts to turn business questions into clear SQL datasets and BI-ready tables, making reporting faster and more consistent.

Tech Stack: Python (pandas/Polars), SQL, Apache Airflow, dbt, PostgreSQL, S3/MinIO, Parquet, Docker, Metabase, BaseLinker, GitHub

EDUCATION

2022 - 2026 | Applied Informatics, Bachelor of Science

Jan Evangelista Purkyně University in Ústí nad Labem (UJEP), Ústí nad Labem, Czechia - Expected May 2026

LANGUAGES

English - B2 (Upper-Intermediate) | Czech - B2 (Upper-Intermediate) | Russian - Native